

## 0.1 Introduction

The framework rests on a single model of the systems it analyses. This chapter defines that model and the object built from it, which every later chapter reads. Section 0.2 describes how a process system is mapped to a directed acyclic graph and discusses the acyclicity assumption. Section 0.3 gives the physical reading of the node roles, and Section 0.4 states the definitions and notation used throughout the thesis. Section 0.5 establishes the layered processing order and the ancestry closures on which every algorithm in the framework depends. Section 0.6 describes the values that an analysis attaches to the network, the three readings of those values, and the three forms in which a value can be known. Section 0.7 describes how a model is built from the files in which real systems are described, and Section 0.8 states the assumptions that all later results are conditional on.

## 0.2 From Process to Graph

A component is any unit of the system that can carry a value of its own, such as a pump or a treatment stage. A dependency is the relation in which one component feeds, supplies or precedes another. The model represents each component as a node and each dependency as a directed edge. In a goal-oriented process system, one in which a resource moves through the components towards a delivery point, the relation between two connected components is already directed. For example, water leaves the intake and arrives at the pump, and the pump does not supply the intake. Hence the direction of an edge is read off the system.

The model further requires that the directed edges form no cycle. For a large class of the systems of interest this is simply true. For example, treatment trains, supply chains, assembly processes and project schedules move material or work strictly forward, and a dependency cycle in a schedule would be a planning error. For other systems, however, acyclicity is an orientation of a richer reality. A distribution pipe can in principle carry water in either direction, but a network operating around its design flows has a dominant direction on every pipe, and orienting each edge by that dominant flow gives a directed acyclic snapshot of the operating regime. What the orientation gives up is the analysis of regimes far from the snapshot, such as flow reversal under unusual demand, recirculation loops and feedback control. The model excludes these by design, and Section 0.8 returns to what that exclusion means for the results. Figure 1 shows the mapping on a small treatment process. The process duplicates its treatment stage, and the duplication appears in the graph as a pair of parallel routes that later reconverge, which is the structural pattern Section 0.3 names.




Figure 1: A small treatment process and its model. Components become nodes, dependencies become directed edges, and the duplicated treatment stage becomes a pair of parallel routes that later reconverge.

## 0.3 Node Roles

The position of a node in the network gives it a role, and each role has a physical reading. A node with no incoming edge is a source, the point where supply, material or work enters the system (e.g. an intake, a generator or the start of a project). Similarly, a node with no outgoing edge is a sink, where the system delivers (e.g. a demand point, an outlet or a completed handover). A node that feeds two or more downstream components is a fork. Forks are how a system distributes and, in particular, how it builds in redundancy, because providing two routes to the same destination begins with one component supplying two. Finally, a node fed by two or more components is a join. This is where distributed flows gather, where assembled parts meet, and where redundant routes come back together.

The last of these readings is a fact about engineered systems as well as about graphs. A robust system duplicates its critical routes by design, and duplicated routes must reconverge at the destination they protect. A well-designed network therefore contains many pairs of paths that separate at a fork and meet again at a join while sharing the components upstream of the fork. Whether that sharing matters depends on the question being asked. Chapter 4 makes the pattern precise, and Chapter 5 shows that for the probabilistic reading of the model, ignoring it gives wrong answers. For the present chapter it is enough to note that reconvergence is a designed feature of the systems being modelled. Figure 2 shows it on a small network.

## 0.4 The Formal Object

**Definition 0.1** (Network model). *A system is modelled as a finite directed acyclic graph  $G = (V, E)$  with node set  $V$  and edge set  $E \subseteq V \times V$ . The sources  $S \subseteq V$  are the nodes with no incoming edge and the sinks  $T \subseteq V$  the nodes with no outgoing edge. For a node  $v$ ,  $\text{Pa}(v)$  denotes its parents, the nodes with an edge into  $v$ , and  $\text{Ch}(v)$  its children.*

**Definition 0.2** (Ancestry). *A node  $u$  is an ancestor of  $v$  when a directed path leads from  $u$  to  $v$ , and a descendant when one leads from  $v$  to  $u$ . The sets  $\text{anc}(v)$  and  $\text{desc}(v)$  collect them. Throughout the framework both are taken self-inclusively, so that  $v \in \text{anc}(v)$  and  $v \in \text{desc}(v)$ .*

**Definition 0.3** (Forks and joins). *The forks of  $G$  are  $F = \{v : |\text{Ch}(v)| \geq 2\}$  and the joins are  $J = \{v : |\text{Pa}(v)| \geq 2\}$ .*

These definitions are the shared vocabulary of the thesis. The chapters that follow refer to them without restating them, and Table 1 collects the notation in one place.




Figure 2: Node roles on a small network. Double borders mark the source and the sink, the bold outline marks the fork, and the shaded node is the join where the redundant pair of routes reconverges alongside an independent supply.

Table 1: Notation established by the network model.

| Symbol                          | Meaning                                                       |
|---------------------------------|---------------------------------------------------------------|
| $G = (V, E)$                    | the system as a finite DAG                                    |
| $S, T$                          | sources (no incoming edge), sinks (no outgoing edge)          |
| $\text{Pa}(v), \text{Ch}(v)$    | parents and children of $v$                                   |
| $\text{anc}(v), \text{desc}(v)$ | ancestors and descendants, self-inclusive                     |
| $F, J$                          | fork nodes (fan-out $\geq 2$ ), join nodes (fan-in $\geq 2$ ) |
| $L = \{L_1, \dots, L_k\}$       | layered order (Section 0.5)                                   |

## 0.5 Layered Order and Closures

Every analysis in the framework visits the nodes of the network in an order in which each node is processed after all of its parents. For a DAG such an order always exists. A topological sort of  $G$  is a linear ordering of  $V$  in which  $u$  precedes  $v$  for every edge  $(u, v) \in E$ . The framework uses a coarser form of the same order, in which nodes at the same depth from the sources are grouped into a layer.

**Theorem 0.4** (Layered order). *For every finite DAG  $G = (V, E)$  there exists a partition of  $V$  into non-empty layers  $L = \{L_1, L_2, \dots, L_k\}$  such that, for every node  $v \in L_i$ , every parent of  $v$  lies in a layer  $L_j$  with  $j < i$ . An algorithm that processes the layers in order, and within a layer processes each node once, visits every node exactly once after all of its parents have been visited, and completes in exactly  $k$  layer steps.*

*Proof.* By induction on  $|V|$ . If  $|V| = 1$  the single node is a source and  $L = \{L_1\}$  with  $L_1 = V$ . Suppose the statement holds for every DAG with fewer than  $n$  nodes, and let  $G$  have  $n$  nodes. Since  $G$  is acyclic it has at least one source; let  $L_1$  be the set of all its sources. Removing  $L_1$  and the edges leaving it gives a DAG  $G'$  with fewer than  $n$  nodes, which by the inductive hypothesis has a layered order  $\{L_2, \dots, L_k\}$  in which every parent within  $G'$  of a node in  $L_i$  lies in an earlier layer. A node of  $G'$  may also have parents in  $L_1$ , and these lie in the earliest layer. Prepending  $L_1$  therefore gives a layered order of  $G$ . The processing claim follows directly: when layer  $L_i$  is processed, every parent of every node in  $L_i$  lies in a layer already processed, so each node is visited once with its parents resolved, and there are  $k$  layers.  $\square$

The theorem matters for the framework in two ways. First, every analysis in the following chapters is an exact method with no iterative convergence loop. Hence a fixed processing order declared once serves all of them, and the number of steps is known from the structure before any analysis runs. Second, nodes within a layer share no dependency, so a layer is a natural boundary for parallel processing by any algorithm that supports it. The layers are computed by a depth-partitioned form of Kahn's algorithm, shown as Algorithm 1, which takes the sources as the first layer, removes them, and takes the nodes that become sources as the next. In the implementation the layers are called iteration sets.

Alongside the layers, later algorithms rely on cheap answers to the question of whether one node is an ancestor of another. The adjacency indices alone do not provide this, since a traversal per query would cost  $O(|V| + |E|)$  each time. Algorithm 1 therefore accumulates the self-inclusive ancestor and descendant closure of every node during the same pass. As each node is dequeued, its ancestor set is passed to its children and it is recorded as an ancestor of each of them, and the descendant sets are populated in the same step. Computing the layers alone costs  $O(|V| + |E|)$ . Carrying the closures adds

a set-union cost that depends on the transitive density of the graph, and in the worst case each node's closure contains every other node, giving  $O(|V|^2)$  space. In return, any reachability query reduces to a single set membership test in constant time. This is a trade of space for time, justified by the number of such queries the decomposition and propagation algorithms make.

---

**Algorithm 1** Layered order with ancestor and descendant closure

---

**Require:** DAG as edge list  $E$  (acyclicity validated by the parser), outgoing index  $\text{Out}$ , incoming index  $\text{In}$

**Ensure:** layers  $L$ , anc, desc

```

1: compute  $\text{indeg}[v]$  for every node  $v$ 
2: initialise  $\text{anc}[v] \leftarrow \{v\}$  and  $\text{desc}[v] \leftarrow \{v\}$  for every node  $v$ 
3: enqueue every node with  $\text{indeg} = 0$  into queue  $Q$ ;  $L \leftarrow \langle \rangle$ 
4: while  $Q$  not empty do
5:    $\ell \leftarrow |Q|$   $\triangleright$  the current layer is the current queue
6:    $L_{\text{cur}} \leftarrow \emptyset$ 
7:   for  $i = 1$  to  $\ell$  do
8:      $x \leftarrow \text{dequeue}(Q)$ ; add  $x$  to  $L_{\text{cur}}$ 
9:     for all  $y \in \text{Out}[x]$  do
10:       $\text{anc}[y] \leftarrow \text{anc}[y] \cup \text{anc}[x]$ 
11:      add  $y$  to  $\text{desc}[a]$  for every  $a \in \text{anc}[x]$ 
12:       $\text{indeg}[y] \leftarrow \text{indeg}[y] - 1$ 
13:      if  $\text{indeg}[y] = 0$  then enqueue  $y$  into  $Q$ 
14:      end if
15:    end for
16:  end for
17:  append  $L_{\text{cur}}$  to  $L$ 
18: end while
19: return  $L$ , anc, desc

```

---

## 0.6 Information on the Network

The model separates what the network is from what is asked of it. The topology of Definition 0.1 is fixed once per system, and an analysis attaches values to its nodes and edges whose meaning depends on the question. Under the probabilistic reading a node's value is the probability that the component operates and an edge's value the probability that the handover between its endpoints succeeds, and the question is with what probability each node is reached from the sources despite component failures. This question is answered by the Probability Propagation Toolkit of Chapter 5. Under the capacity reading the values are throughput limits, and the question is how much the network can deliver, where the limits bind and what relieving them would gain. This question is answered by the Capacity Flow Toolkit of Chapter 6. Under the schedule and cost reading the values are durations or costs, of the component's own processing and

of the transfer along each dependency, and the question is when the network completes, which dependencies set that outcome and which have room to move. This question is answered by the Critical Path Toolkit of Chapter 7. The three readings are three different analyses of one topology. Because the topology is common to them, a single system model answers questions that are usually asked of three separate models [1].

How well a value is known is a further property, independent of what it means. Chapter 2 defined the three forms the framework accepts. A deterministic value asserts full knowledge and is appropriate for design constants and well-measured components. An interval asserts only bounds and is appropriate when data is sparse, when a manufacturer quotes a tolerance, or when experts will commit to a range but not to a number. A probability box bounds a distribution and carries partial statistical knowledge, for example measurements sufficient to constrain a curve without fixing it. The framework propagates each form through the analyses that support it, as set out in Chapter 1, so that the choice among them can be made on the state of knowledge about the system.

Before any of the three questions can be asked, the model must be built from the descriptions in which real systems arrive, which is the subject of the next section, and for the probabilistic reading the reconvergent structure of Section 0.3 must be identified, which is the subject of Chapter 4. Both are preprocessing steps performed once per system.

## 0.7 From File to Model

The data describing an infrastructure system is rarely held in one place or in one format. Different objectives are pursued on the same system, a fault-detection workflow tracks health signals while a planning workflow tracks costs, and each maintains the parameters relevant to its own purpose. Furthermore, the tools used across these workflows implement different schemas and export conventions, and a dataset is often shaped to fit a particular tool. Beneath these differences, however, any DAG description of a system reduces to a node list and an edge list, in which nodes are point features such as sources, junctions, assets or demand points and edges are the relationships between them. Infrastructure formats such as GraphML, EPANET `.inp` and MATPOWER `.m` all reduce to node and edge records beneath their schemas, and the framework asks the user to perform that reduction, supplying the structure in one of the two forms described in Section 0.7.1. Values describing the behaviour of the system, such as a capacity or a failure probability, are a separate layer of information over that structure, and a source dataset may combine the two. The input module treats structure and analysis inputs as two layers of one model, parses each from its own file, and produces the graph object every analysis reads. It is the single mandatory entry point of the framework: no analysis can be invoked without it.

### 0.7.1 Network structure files

Components in a physical system usually have descriptive identifiers, such as `Pump_A` or `Substation_North`. However, the framework requires integer identifiers, because integers are compact to store and fast to index in the repeated lookups and set operations that every algorithm performs. The integers need not be contiguous, but the user is responsible for mapping domain identifiers to unique integers before parsing.

The network parser accepts an edge list or an adjacency matrix, in a `.edges` or a `.csv` file, and detects which it has been given in stages. A `.edges` extension selects the edge-list reader directly. A `.csv` file is ambiguous, so the parser first scans the header for source and destination keywords that indicate an edge list; failing that, it scans the content for the integer, delimiter, integer pattern of an edge record; failing that, it treats a two-column table as an edge list and anything else as an adjacency matrix. An adjacency matrix is streamed row by row into an edge list. Header rows and comment lines are skipped in either case.

As edges are read, self-loops are rejected at once and cycle-forming edges are detected during ingestion, since both violate the acyclicity of Definition 0.1. For every valid edge both nodes and both directions of the relationship are recorded in the same pass, so that the incoming and outgoing indices are built together and no second pass over the edges is needed to construct the reverse adjacency. Ingestion is therefore linear in the number of edges, and the parser returns the minimal graph: the edge list, the two indices and the sources.

### 0.7.2 Analysis input files

The values that drive an analysis are usually observed and managed separately from the network structure, in manufacturer specifications, historic records or planning estimates, and the framework keeps them in separate files. Embedding them in the structure file would produce a file that is large and hard to maintain. It would also force a user who wants to evaluate the same network under several readings, or under several scenarios of one reading, to keep a duplicate topology file for each.

All analysis inputs are given as JSON files. JSON is human-readable, platform-independent, and maps naturally onto the key-value form in which component attributes are described. Where a flat format such as CSV suits the rigid tabular shape of an edge list, JSON's nesting suits the values themselves. An interval or a p-box can be written directly as a nested object without inventing column conventions or flattening a structured value into several cells. No specialist tool is needed to write or edit such a file, and a data owner who understands the node and edge identifiers can contribute values without knowledge of graph theory.

Every analysis input file declares a `data_type`, which is `Float64`, `Interval` or a

probability-box construction, and the framework selects the deserialiser from that declaration. The analysis class is not declared in the file; it is fixed by the file’s role, since each analysis reads its own input file from the session. Every key must map to a valid node identifier or edge identifier of the parsed graph, and every value must be parseable into the declared type. The parser is driven by the values and the declared type. Descriptive labels are ignored, so the JSON can carry whatever naming convention the data owner uses. Concentrating the deserialisation of uncertain values in one place, instead of in a separate reader per analysis, keeps the handling of intervals and p-boxes consistent across the framework.

The three analysis classes differ in what they accept. Probability inputs are given in two files, one carrying a prior for each node and one carrying a probability for each edge. All values must lie in  $[0, 1]$ , and values outside that range are rejected at parse time. This is the most expressive class and accepts all three data types. An interval is written as a lower and an upper value. When the file declares the `pbox` data type, each value is an object that names its own construction, so that values in one file may use different constructions. The supported constructions mirror the constructor families of the `ProbabilityBoundsAnalysis.jl` library: a scalar or an interval as a degenerate box, a parametric distribution with point or interval parameters, an envelope of distributions, or a distribution-free box from bounds on the moments. Flow capacity inputs are given in one file as a list of edge records, each carrying a source, a destination and a capacity, with an optional list of node capacities. Since capacity is point-valued in the present framework, only the `Float64` data type is accepted. Unconstrained capacity is written with an infinity token, and every capacity is checked to be non-negative and not NaN at parse time. Critical path inputs are given in one file with a time section carrying a duration for each node, a delay for each edge and an initial time, and an optional cost section of the same shape. This class accepts `Float64` and `Interval`. The framework is unit-agnostic for it: the values may be durations, costs or any other quantity that accumulates along a path, provided the interpretation is consistent within a run, and the analysis mode that reads them is chosen at invocation as described in Chapter 7. Table 2 summarises the contracts.

### 0.7.3 The graph object

The output of input processing is the framework’s graph object, assembled once from the validated outputs of the parser and of Algorithm 1. It carries the edge list, the outgoing and incoming indices, the sources, the sinks, the sorted node list, the fork and join sets, the layers, and the ancestor and descendant closures. These fields are built once and shared by every analysis. The analysis inputs, node priors and edge probabilities, edge and node capacities, and node and edge values, are attached to the same object per

Table 2: Input contracts by layer and analysis class.

| Layer or class    | Contract                                                                        | Data types                               | Validation                                                     |
|-------------------|---------------------------------------------------------------------------------|------------------------------------------|----------------------------------------------------------------|
| Network structure | Edge list or adjacency matrix over integer node identifiers                     | Integers; directed edges                 | Format detected in stages; self-loops and cycles rejected      |
| Probability       | Node prior file and edge probability file, keyed to graph identifiers           | <b>Float64</b> , <b>Interval</b> , p-box | Values in $[0, 1]$ ; <b>data_type</b> selects the deserialiser |
| Flow capacity     | Edge records (source, destination, capacity); optional node capacities          | <b>Float64</b> only                      | Non-negative and not NaN; infinity token accepted              |
| Critical path     | Time section (node durations, edge delays, initial time); optional cost section | <b>Float64</b> , <b>Interval</b>         | Unit-agnostic; mode chosen at invocation                       |

invocation, so that one instance of a network can be evaluated under different readings and under different scenarios of one reading without rebuilding the structure. Every argument that any analysis in the framework requires can be taken from this object, and it is the point of handover from preprocessing to analysis. The concrete types chosen for its fields, and the reasons for building a purpose-specific object instead of using the language’s general graph library, are discussed in Chapter 8.

## 0.8 Assumptions and Limits

The model rests on four assumptions, and each is a limit on what the results mean.

1. Component behaviours are independent. A common cause that fails several components together is not represented, and where it matters it must be modelled explicitly as a shared upstream node.
2. Under the probabilistic reading a component is binary, operational or failed, and degraded intermediate states are outside the model.
3. The topology is static over the analysis horizon, so reconfiguration during operation is treated as a sequence of separate analyses.
4. The acyclicity of Section 0.2 excludes feedback and flow reversal. This is a property of the operating snapshot being analysed, since the physical assets may carry flow in either direction.

Every analysis in the following chapters assumes this list, and none of the algorithms introduces an assumption beyond it. Where a later chapter relaxes or adds to the list for a particular reading, it says so.

## 0.9 Summary

A process system is modelled as a finite directed acyclic graph in which components are nodes, dependencies are directed edges, and the roles of source, sink, fork and join have physical readings. The graph object built from the model carries the topology, a layered processing order whose existence Theorem 0.4 guarantees, and the ancestor and descendant closure of every node, all computed in one pass. An analysis attaches values to the object under one of three readings, probabilistic, capacity, or schedule and cost, in one of three forms, deterministic, interval or probability box. The input module builds the object from edge-list or adjacency-matrix structure files and JSON analysis files whose contracts were given in Section 0.7, validating acyclicity, identifier mapping and value domains at parse time. The four assumptions of Section 0.8 apply to every later result.

# Bibliography

- [1] M. Ouyang, “Review on modeling and simulation of interdependent critical infrastructure systems,” *Reliability Engineering & System Safety*, vol. 121, pp. 43–60, 2014. DOI: 10.1016/j.ress.2013.06.040