

Information Propagation Algorithm for Exact Reliability Analysis in Directed Acyclic Process Networks

Temitope Ohiani^a, Edoardo Patelli^a, Caroline Pyke^a

^a *University of Strathclyde, Glasgow, United Kingdom*

Abstract

Reliability analysis in directed acyclic process networks requires computing, for each node, the probability that at least one viable path exists from any source node to that node. Standard analytical formulations become inaccurate or inefficient in the presence of reconvergent path structures, where multiple paths diverge from shared upstream ancestors and reconverge downstream, creating statistical dependencies that violate path-independence assumptions. This paper presents the Information Propagation Algorithm (IPA), an exact message-passing method that resolves these dependencies through conditional enumeration on shared fork ancestors. IPA processes the network in topological order, identifying reconvergent (diamond) substructures, solving them by conditioning on a minimal separating set, and replacing them with statistically equivalent supernodes. The method is shown to be a specialisation of cutset conditioning to source-to-node reachability: its per-instance cost is exactly determined by the conditioning-set sizes encountered, a width parameter that also governs the size of a well-ordered binary decision diagram on the same network. Exactness is verified against an independent reduced-ordered binary decision diagram across a corpus of 129 random directed acyclic graphs, six topological families, and several real infrastructure networks, with worst observed disagreement at floating-point precision. The distinguishing capability of the method is the native propagation of imprecise component reliabilities: interval-valued inputs yield the exact reliability range, a consequence of the monotonicity of the reachability function, and probability-box inputs yield guaranteed distributional bounds through a conditioning operator built on the Fréchet–Hoeffding inequalities, enabling certified bounds on decision-relevant probabilities that point-valued

exact methods cannot produce analytically. The method is applied to a medical drone logistics network for Scotland in which every input is derived from the source design study. Propagated reliability bounds identify the locations whose delivery reachability is most uncertain, and a sweep of a network-redundancy design parameter locates the practical boundary of exact computation on this real network; within the range tested, the proposed method remains practical at a redundancy level at which the decision-diagram alternative did not complete within a practical computational budget.

Keywords: Network Reliability, Directed Acyclic Graphs, Message Passing, Reconvergent Structures, Imprecise Probability, Process Networks, Infrastructure Systems

1. Introduction

Many infrastructure and engineered systems operate through the propagation of resources, services, or information across interconnected components. Examples include electric power networks, water distribution systems, telecommunication networks, transportation systems, and logistics networks. In such systems, functionality depends not only on the reliability of individual components but also on the existence of viable directed paths through which flow can propagate from source nodes to downstream nodes. In this paper, we refer to this class of systems as multi-source process networks. Failures in these systems can have severe operational, economic, environmental, and societal consequences. Their assessment therefore requires metrics that quantify not only component reliability but also the ability of the network to preserve connectivity and sustain flow. A natural reliability measure in this context is the probability that at least one viable path exists from one or more source nodes to a given target node. This source-to-node reachability reliability provides a direct measure of the network's ability to maintain function under component and transmission failures.

Directed acyclic graphs (DAGs) are a natural fit to model and evaluate the performance of resource flow systems. In a DAG, nodes represent system components, processing stages, or transfer locations, while directed edges model permissible flow transitions and dependencies between components. Consider the typical process of transport networks where service at downstream locations may depend on the successful operation of several upstream links, hubs, or transfer stages. Failures at earlier stations can cause

disruption and delays in travel and, in some cases, stop service altogether.

To mitigate cascading failure scenarios and service bottlenecks, the reliability of the network’s resource flow is an important performance metric to monitor. Reliability is generally defined as the probability that a system performs its intended function over a specified time. In flow networks, the intended function is that the other parts of the network successfully receive information flowing from the sources. Complex systems have redundancy strategies and backup paths to ensure this flow. Reliability in multi-source flow networks is related to the probability that at least one viable operational path exists from source nodes to the other nodes in the network.

The classic approach to exactly solving the reliability problem involves enumerating and combining all possible source-to-node paths. However, this quickly becomes computationally intractable for complex process systems with multiple sources, extensive redundancy, and re-convergent network topologies (George-Williams et al. 2021). In particular, when distinct paths share upstream nodes or edges before reconverging downstream, the resulting path events are no longer independent, and naïve path-combination rules become invalid.

The structure of DAGs suggests an alternative viewpoint based on message passing. Rather than globally enumerating all viable paths, reliability can be computed locally and recursively by asking, at each node: what is the probability that at least one viable signal reaches this node from any source? For tree-like or conditionally independent substructures, this can be evaluated efficiently using local update rules. However, reconvergent path structures remain problematic because they induce statistical dependencies between incoming signals. These dependencies must be resolved exactly if message passing is to remain accurate.

This paper presents the Information Propagation Algorithm (IPA), an exact message-passing method for reliability analysis in directed acyclic process networks. IPA exploits topological ordering to propagate source-to-node reliability from upstream to downstream nodes. Its central idea is to identify reconvergent fork–join subgraphs, referred to here as diamond structures, and resolve their dependency explicitly by conditioning on the shared fork ancestors. Once solved, these diamond subgraphs are replaced with statistically equivalent supernodes, allowing the network to be progressively simplified without loss of exactness.

The proposed approach presents several advantages. First, it preserves exact reliability computation in DAGs with reconvergent structures, avoiding

the independence errors that arise from standard local update rules. Second, by decomposing and caching repeated dependency patterns, it avoids repeated conditional enumeration over already-resolved substructures, and its computational cost on a given network is determined exactly, before enumeration begins, by the sizes of the conditioning sets its decomposition identifies. Third, and distinctively, the propagation scheme extends beyond point-valued component reliabilities: interval-valued reliabilities are propagated exactly, and probability-box (p-box) reliabilities are propagated with guaranteed bounds, so that epistemic uncertainty about component data can be carried through the analysis rather than suppressed into point estimates. These properties make the method suitable for design-stage studies in which many network configurations must be assessed, and for decision problems in which the reliability inputs themselves are only partially known.

1.1. Related Works

Early reliability analysis methods are based on relatively simple structural representations of systems. The Reliability Block Diagram (RBD) (Billinton and Allan 1992; Rausand 2014; Ebeling 2019) method models a system as interconnected component blocks arranged in series, parallel, or k -out-of- n configurations, whose reliability can be computed using closed-form expressions. While computationally efficient and intuitive, RBDs are limited in their ability to represent complex dependency structures and reconvergent flows. Extensions such as Dynamic Reliability Block Diagrams (DRBD) (Distefano and Puliafito 2007) incorporate time-dependent behaviour, allowing for more realistic modelling of component interactions and system dependencies, but they still struggle to represent highly interconnected networks with multiple convergent paths.

Boolean logic-based methods provide a more flexible framework for modelling system failure mechanisms by including event causes and consequences. Fault Tree Analysis (FTA) (Xing and Amari 2008) is a top-down deductive logic tree to identify possible combinations of component failures that could cause specific undesirable events. Event Tree Analysis (ETA) is the complementary inductive method where the logic tree instead begins with the initiating event and systematically traces all possible sequences of subsequent successes and failures through branching structures and logic gates (Čepin 2011). These methods can capture a wider range of logical dependencies than RBDs; however, for large-scale systems they require extensive manual construction and become difficult to maintain. In addition, their

computational cost grows rapidly with the size and complexity of the logical expressions, particularly in systems with multiple interacting pathways (Sherwin and Bossche 1993; Xing and Amari 2008).

The main limitations of tree-based methods stem from exponential growth in logic expressions and manual construction requirements for complex systems. Binary Decision Diagrams (BDDs) (Rauzy 1993; Bryant 1986; Brace et al. 1990) were adapted from circuit design to improve fault tree models by using variable elimination and path reduction techniques on Boolean functions. In BDDs, Boolean functions are represented as directed acyclic graphs based on Shannon decomposition, where each non-sink node represents a Boolean variable with two outgoing edges corresponding to the variable being true or false (Xing and Amari 2008). This structure allows BDD-based methods to require less memory and computational time than traditional fault tree analysis (Rauzy 1993). BDD algorithms can handle complex Boolean expressions with thousands of variables while maintaining polynomial space complexity for many systems. However, their performance depends critically on variable ordering, and determining an optimal ordering is itself computationally intractable in general (Bryant and Meinel 2002). For highly interconnected networks with shared components and reconvergent fanout structures, BDD construction with diamond structures — where multiple independent paths converge at common nodes — suffers from the reconvergence problem, requiring extensive intermediate node generation that can exceed memory limitations even for moderately sized systems. These computational challenges become particularly severe when modelling common cause failures and complex dependency structures that propagate through multiple interconnected pathways simultaneously.

Common cause failures create statistical dependencies that violate the independence assumptions underlying traditional reliability calculations (Mosleh 1993). When multiple components share common failure modes, the resulting dependency patterns cannot be efficiently captured through simple Boolean combinations, representing a fundamental limitation shared by RBDs, FTAs, and BDDs in process systems with convergent network topologies.

To overcome these limitations, path-based and graph-theoretic methods have been developed. These methods contextualise the network as a probabilistic graph theory problem and leverage probability mathematics and set operations to calculate network reliability. Exact methods based on minimal path sets and minimal cut sets (Kumar et al. 2018; Li and Zhao 2005; Billinton and Allan 1992; Liu and Li 2012) identify the smallest collections

of components that ensure connectivity or guarantee system failure. Recursive decomposition techniques further improve efficiency by partitioning the network into smaller subproblems. For example, the cut-based method in Liu and Li (2012) proposes simultaneously using both the disjoint minimal cut sets and disjoint minimal path sets during the decomposition process for networks with low edge reliabilities. Kim and Kang’s method (Kim and Kang 2013) extends the Recursive Decomposition Algorithm through a Logical Expansion approach (LE-RDA) that uses two separate procedures for intersection and union logical functions to handle decomposition of general system events with multiple node pairs. Nevertheless, these methods remain computationally expensive for networks with extensive redundancy and re-convergent connectivity, as they still require combinatorial enumeration of path or cut structures.

Simulation methods such as Monte Carlo (MC) have become a widely adopted reliability analysis method for complex networks (Fishman 1986). MC algorithms can handle arbitrary network structures and provide confidence intervals, although they can become computationally expensive for large networks requiring many samples to reach acceptable accuracy. Monte Carlo variants such as importance sampling (Hui 2007) and subset simulation (Au and Beck 2001) have been proposed to address rare failure events, though these optimisations often demand more computational cost.

More recently, message passing approaches based on probabilistic graphical models have been explored for network reliability analysis (Yedidia et al. 2003; Pearl 1988). These methods infer the reliability of the system by exploiting the network structure to propagate the local calculations throughout the network until convergence. Unlike enumeration-based methods, message passing scales linearly with network size and naturally handles complex dependency structures. For instance, methods for system reliability such as the directed Probability Propagation Method (dPrPm) (Tong and Tien 2019) also allow dependencies to be handled through pairwise probability tracking and approximation of joint distributions. While effective, these approaches can become memory-intensive and computationally demanding as the number of dependent interactions increases. In addition, convergence is not always guaranteed (Murphy et al. 2013), and approximation quality depends heavily on network topology and component correlation strength.

In the exact-inference literature on probabilistic graphical models, dependencies induced by shared ancestry are classically resolved either by conditioning — cutset conditioning enumerates the states of a set of variables

whose removal renders the remaining structure singly connected (Pearl 1988) — or by clustering, as in the junction-tree algorithm, whose cost is exponential in the treewidth of the network (Lauritzen and Spiegelhalter 1988). Both families are exact and both are governed by a structural width parameter of the underlying graph. Section 4.3 positions the present method within this landscape: it is a specialisation of cutset conditioning to the source-to-node reachability problem, and its realised cost is governed by the same width parameter that controls junction-tree inference and well-ordered binary decision diagrams.

A further limitation cuts across all of the exact methods above: they operate on point-valued component reliabilities. In practice, component failure probabilities are often known only imprecisely — from sparse data, expert elicitation, or manufacturer bounds — and a rich literature on imprecise probability represents such knowledge as intervals or as probability boxes (p-boxes), i.e. bounds on the cumulative distribution of an uncertain quantity (Ferson et al. 2003). Arithmetic on p-boxes with known or unknown dependence is formalised in Williamson and Downs (1990), building on the classical Fréchet–Hoeffding bounds on joint distributions. Decision-diagram and path-based exact methods do not propagate such imprecise inputs natively: an interval query requires repeated evaluation at selected input corners, and no analytic route exists to a distributional bound. The method proposed here propagates interval and p-box reliabilities directly, which is the capability that distinguishes it from a well-ordered decision diagram of comparable computational cost.

Hence, existing methods face a trade-off between generality, computational efficiency, and exactness. Classical analytical methods struggle with complex dependency structures, exact path-based approaches suffer from combinatorial explosion, and simulation or approximate message-passing methods sacrifice exactness for scalability. This work addresses this gap by proposing an exact message-passing framework for DAG reliability analysis that explicitly resolves reconvergent dependencies through conditional decomposition of fork–join substructures. By identifying and reducing these structures during computation, the proposed method avoids repeated enumeration over shared dependencies while preserving exact probabilistic results.

2. Network Model

A process system is modelled as a directed acyclic graph $G = (V, E)$ where V denotes the set of nodes (components) and $E \subseteq V \times V$ represents the set of directed edges (dependencies). Let $S \subseteq V$ denote the set of source and $T \subseteq V$ the sink nodes. Throughout the paper we use $\text{Pa}(v)$ and $\text{Ch}(v)$ to represent the set of parent nodes and child nodes of v , respectively.

2.1. Component States and Probabilistic Model

Each node $v \in V$ is associated with a binary random variable $X_v \in \{0, 1\}$, where $X_v = 1$ indicates that node v is operational. Similarly, each directed edge $(u, v) \in E$ is associated with a binary random variable $X_{uv} \in \{0, 1\}$ representing successful transmission from u to v . We define:

- $R(v) = P(X_v = 1)$: intrinsic reliability of node v ;
- $P(u, v) = P(X_{uv} = 1)$: transmission probability of edge (u, v) .

A directed path $\pi(s, v) = (s = v_0, v_1, \dots, v_k = v)$ is said to be viable if all nodes and edges along the path are operational, i.e.

$$X_{v_0} X_{v_1} \cdots X_{v_k} \cdot X_{v_0 v_1} \cdots X_{v_{k-1} v_k} = 1. \quad (1)$$

Definition 1 (Viable Path). A path $\pi(s, v)$ from source $s \in S$ to node $v \in V$ is viable if all nodes and edges along the path are functional.

2.2. Definition of Diamond Structures

We formalise reconvergent dependency patterns through the notion of a diamond subgraph.

Definition 2 (Diamond Structure). A diamond structure is a subgraph $D \subseteq G$ characterised by: a set of fork nodes $F \subseteq V$, where each $f \in F$ has at least two outgoing paths; a join node $j \in V$ with at least two incoming paths; at least two distinct directed paths from nodes in F to j ; and at least one shared upstream node or edge among these paths.

In such structures, the paths from fork nodes to the join node are not independent because they share common components. The simplest case is a single-fork diamond, where all reconvergent paths originate from a single fork node.




Figure 1: A multi-source directed acyclic graph.

2.3. Reliability Objective

The reliability problem is formally stated as computing the probability that there is at least one viable path from any source node to each target node in the network. The reachability probability of a node $v \in V$ is defined as the probability that at least one viable path exists from any source node to v :

$$b(v) = P\left(\bigcup_{s \in S} \bigcup_{\pi \in \Pi(s,v)} \{\pi \text{ is viable}\}\right), \quad (2)$$

where $\Pi(s, v)$ denotes the set of all directed paths from s to v . The quantity $b(v)$ represents the probability that node v successfully receives flow from at least one source and is referred to as the node reliability (or belief) at v .

Equivalently, reachability can be expressed recursively. Define for each node the reachability indicator

$$\mathcal{R}_v = X_v \wedge \left(v \in S \vee \bigvee_{u \in \text{Pa}(v)} (\mathcal{R}_u \wedge X_{uv}) \right), \quad (3)$$

so that $b(v) = P(\mathcal{R}_v = 1)$. Because all node and edge indicators are mutually independent (Assumption 2), $\mathcal{R}_v$ is a monotone Boolean function of the component indicators: switching any component from failed to operational can only preserve or create source-to-node connectivity, never destroy it. Two consequences are used repeatedly in what follows. First, computing $b(v)$ is

#P-hard in general, since it subsumes two-terminal network reliability. Second, the monotonicity of $\mathcal{R}_v$ in every component makes $b(v)$ a non-decreasing function of every component reliability, a property that Section 5.3 exploits to propagate interval-valued inputs exactly.

For a node v , let $\text{anc}(v)$ denote the ancestors of v (nodes with a directed path to v), and let $F = \{v : |\text{Ch}(v)| \geq 2\}$ and $J = \{v : |\text{Pa}(v)| \geq 2\}$ denote the fork and join nodes of G . Given a set $E_c \subseteq F$ of forks whose reachability states have been fixed by conditioning, the *unresolved influence* of a parent u of a join node is

$$\text{infl}(u; E_c) = (\{u\} \cup \text{anc}(u)) \setminus E_c, \quad (4)$$

the upstream components on which the signal through u still depends. Two parents are dependent precisely when their unresolved influences intersect; this is the structural criterion used by the decomposition of Section 4.

2.4. Model Assumptions

Assumption 1 (Graph Structure). The network G is a directed acyclic graph (i.e. G contains no directed cycles), and every node $v \notin S$ is reachable from at least one source node.

Assumption 2 (Component Independence). Node states $\{X_v\}$ and edge states $\{X_{uv}\}$ are mutually independent. Node and edge states are also independent of each other.

Assumption 3 (Binary Operation). All nodes and edges have binary states (operational or failed).

Dependencies between paths arise solely from shared nodes or edges. In particular, signals arriving at a node from different parents are independent only if the corresponding upstream paths do not share common components. This formulation is equivalent to a network reliability problem with node and edge failures, where the objective is to compute source-to-node reachability probabilities.

2.5. Imprecise Component Reliabilities

The formulation above takes the component reliabilities $R(v)$ and $P(u, v)$ as known point values. In many applications these inputs are themselves uncertain: failure probabilities may be estimated from sparse operational

data, elicited from experts, or specified only as bounds. We therefore allow each component reliability to be specified in one of three forms of increasing informativeness:

1. a point value $p \in [0, 1]$, as above;
2. an interval $[\underline{p}, \bar{p}] \subseteq [0, 1]$, asserting only that the true reliability lies between the stated bounds;
3. a probability box (p-box): a pair of cumulative distribution functions $(\underline{F}, \bar{F})$ bounding the distribution of the uncertain reliability, appropriate when partial distributional information is available (Ferson et al. 2003; Williamson and Downs 1990).

The reliability objective generalises accordingly: for interval inputs, the target is the exact range $[b(v), \bar{b}(v)]$ of node reliabilities attainable over all component reliabilities consistent with the stated intervals; for p-box inputs, the target is a sound pair of bounds on the distribution of $b(v)$. Section 5.3 shows that the proposed propagation scheme computes the interval range exactly — a consequence of the monotonicity established above — and computes guaranteed distributional bounds in the p-box case.

3. The Information Propagation Algorithm

The Information Propagation Algorithm (IPA) computes source-to-node reliability in a directed acyclic graph by propagating reachability probabilities from upstream nodes to downstream nodes in topological order. For each node $v \in V$, the objective is to compute $b(v)$, the probability that at least one viable path reaches v from any source.

3.1. Signal Propagation Formulation

Let v be a non-source node with parent set $\text{Pa}(v) = \{p_1, p_2, \dots, p_k\}$. For each parent $p_i \in \text{Pa}(v)$, define the incoming signal event $A_i = \{\text{a viable signal reaches } v \text{ through } p_i\}$. The corresponding signal probability is

$$S_i = P(A_i) = b(p_i) P(p_i, v), \quad (5)$$

that is, the probability that a viable path reaches p_i and is successfully transmitted through edge (p_i, v) . Since node v must itself be operational in order to receive and transmit flow, its reliability is

$$b(v) = R(v) P\left(\bigcup_{i=1}^k A_i\right). \quad (6)$$




Figure 2: Single-path structure: chained signal propagation.

For a source node $v_s \in S$, reliability is simply its intrinsic reliability: $b(v_s) = R(v_s)$.

3.2. Topological Evaluation Strategy

Because the network is acyclic, node reliabilities can be evaluated in topological order. Let $L = \{L_1, L_2, \dots, L_K\}$ denote a layering of the graph such that every edge goes from an earlier layer to a later one. Then, for any node $v \in L_i$, all parent nodes $p \in \text{Pa}(v)$ belong to layers L_j with $j < i$. This property allows IPA to evaluate all node reliabilities in a single forward pass: source nodes are initialised with their intrinsic reliabilities; each non-source node is processed once all parent reliabilities are available; and the local update depends on whether the incoming parent signals are independent or dependent. No iterative convergence procedure is required.

3.2.1. Independent Parent Signals

If the incoming signal events $\{A_i\}_{i=1}^k$ are mutually independent, then the probability that node v receives at least one signal is

$$P\left(\bigcup_{i=1}^k A_i\right) = 1 - \prod_{i=1}^k (1 - S_i). \quad (7)$$

Substituting (7) into (6) gives

$$b(v) = R(v) \left(1 - \prod_{p \in \text{Pa}(v)} (1 - b(p) P(p, v)) \right). \quad (8)$$

Equation (8) is exact when the upstream paths associated with the incoming parent signals do not share common nodes or edges. Two important special cases are shown in Figures 2 and 3. For a node with a single parent p , the update reduces to the chained-path case $b(v) = R(v) b(p) P(p, v)$.




Figure 3: Multiple independent parent signals: parallel signal propagation.

3.2.2. Reconvergent Structures and Independence Breakdown

The update rule in (8) is no longer exact when the incoming signal events are dependent. In directed acyclic graphs, this occurs in reconvergent structures, where multiple upstream paths share common components before merging at a downstream node. Let v be a join node receiving signals from parents $\{p_1, p_2\}$. If the upstream paths leading to p_1 and p_2 share a common ancestor f , then the events A_1 and A_2 are dependent. In that case, applying the independence-based rule $P(A_1 \cup A_2) = 1 - (1 - S_1)(1 - S_2)$ overestimates the true probability of signal reception, because it ignores the correlation induced by the shared upstream components. IPA resolves this problem by identifying such reconvergent structures and conditioning on the shared upstream variables that induce the dependency.

3.2.3. Diamond Structures and Conditional Decomposition

The fundamental dependency-inducing pattern considered by IPA is a diamond structure: a reconvergent fork–join subgraph in which multiple paths diverge from one or more shared upstream nodes and reconverge at a join node. Consider a single-fork diamond with fork node f and join node j . Let $J = \{\text{a signal reaches } j\}$ and $F = \{\text{fork node } f \text{ is reached and operational}\}$. By the law of total probability,

$$P(J) = P(J | F)P(F) + P(J | \bar{F})P(\bar{F}). \quad (9)$$

Conditioning on the state of the fork restores independence among the downstream paths: if $F = 1$, the fork is active and the paths downstream of f become conditionally independent; if $F = 0$, no signal can propagate through paths originating from f . Thus, if the diamond contains k downstream paths

from f to j , then

$$P(J | F) = 1 - \prod_{i=1}^k (1 - S_i(F)), \quad (10)$$

where $S_i(F)$ denotes the probability that path i transmits a signal to j given that the fork is active. For a simple induced diamond with no alternative upstream routes to the join node, $P(J | \bar{F}) = 0$.

3.2.4. Generalisation to Multiple Conditioning Nodes

More complex reconvergent structures may involve several shared upstream nodes $C = \{c_1, c_2, \dots, c_m\}$. In that case, the join-node signal probability is computed by enumerating all conditioning states:

$$P(J) = \sum_{s \in \{0,1\}^m} P(J | s) P(s), \quad (11)$$

where $s = (s_1, \dots, s_m)$ is a binary assignment to the conditioning nodes and

$$P(s) = \prod_{i=1}^m [b(c_i)]^{s_i} [1 - b(c_i)]^{1-s_i}. \quad (12)$$

For each conditioning state s , the remaining incoming paths to the join node become conditionally independent, so the local update can again be evaluated using the independent-signal formula.

3.3. Simple Example

To illustrate the effect of reconvergent dependencies and the correctness of the conditioning approach, consider the simple diamond in Figure 4. The subgraph consists of a single fork–join pair connected by two intermediate paths. All non-source nodes have intrinsic reliability 0.9, and all edges have transmission probability 0.9.

3.3.1. Solution by Path Enumeration (Correct but Expensive)

We enumerate the relevant fork/intermediate-node states:

$$\begin{aligned} P_1 &= 0.81 \times 0.81 \times 0.81 = 0.531441 && \text{(fork active, both intermediates work),} \\ P_2 &= 0.81 \times 0.81 \times 0.19 = 0.124659 && \text{(fork active, only first intermediate works),} \\ P_3 &= 0.81 \times 0.19 \times 0.81 = 0.124659 && \text{(fork active, only second intermediate works),} \\ P_4 &= 0.81 \times 0.19 \times 0.19 = 0.029241 && \text{(fork active, both intermediates fail),} \\ P_5 &= 0.19 && \text{(fork inactive, no signal possible).} \end{aligned}$$




Figure 4: A simple diamond subgraph between fork node f and join node j .

The signal-reception probabilities at the join node are 0.99 if both paths send signals, 0.9 if only one path sends a signal, and 0 if no path sends a signal. Hence,

$$\begin{aligned} P(\text{signal}) &= P_1 \times 0.99 + P_2 \times 0.9 + P_3 \times 0.9 \\ &= 0.531441 \times 0.99 + 0.124659 \times 0.9 + 0.124659 \times 0.9 = 0.75051279. \end{aligned}$$

The resulting reliability at the join node is $b(j) = R(j) P(\text{signal}) = 0.9 \times 0.75051279 = 0.675461511$.

3.3.2. Solution by Independent Update (Incorrect)

Using the independent-parent update yields

$$\begin{aligned} b(f) &= R(s) P(s, f) R(f) = 1.0 \times 0.9 \times 0.9 = 0.81, \\ b(p_1) &= b(p_2) = b(f) P(f, p) R(p) = 0.81 \times 0.9 \times 0.9 = 0.6561. \end{aligned}$$

Thus $S_1 = S_2 = b(p_i) P(p_i, j) = 0.6561 \times 0.9 = 0.59049$. Treating the two incoming signals as independent gives

$$\begin{aligned} P(\text{signal}) &= 1 - (1 - 0.59049)^2 = 0.83227, \\ b(j) &= 0.9 \times 0.83227 = 0.749043, \end{aligned}$$

which overestimates the correct value.

3.3.3. Solution by Diamond Conditioning (Correct and Efficient)

Conditioning on the fork state, we have $b(f) = 0.81$. If the fork is active, each downstream path has conditional transmission probability $P(f, p) R(p) P(p, j) = 0.9 \times 0.9 \times 0.9 = 0.729$. Therefore,

$$P(\text{signal} \mid F = 1) = 1 - (1 - 0.729)^2 = 0.926559, \quad P(\text{signal} \mid F = 0) = 0.$$

Using (9),

$$P(\text{signal}) = 0.926559 \times 0.81 + 0 \times 0.19 = 0.75051279,$$

hence $b(j) = 0.9 \times 0.75051279 = 0.675461511$, which matches the exact path-enumeration result. This example shows the central idea behind IPA: the dependence induced by a reconvergent structure can be handled exactly by conditioning only on the shared upstream variables that generate the dependence.

3.4. IPA Workflow and Algorithm

The Information Propagation Algorithm computes node reliability in a single forward pass over the DAG. For each node, the algorithm: computes incoming signal probabilities from its parent nodes; applies the closed-form independent update when the parent signals are independent; detects reconvergent structures when dependencies are present; and resolves those dependencies through conditional enumeration on shared upstream nodes. This formulation localises exact dependency handling to the subgraphs that require it, while retaining simple closed-form updates elsewhere. Algorithm 1 summarises the overall IPA procedure.

3.4.1. Computational Characteristics

For nodes with independent parent signals, the update cost is $O(|\text{Pa}(v)|)$, so the total computational cost is linear in the number of edges for tree-structured networks. For nodes involved in reconvergent structures, additional cost arises from conditional enumeration over shared upstream components. If a diamond requires conditioning on c shared upstream variables, the local cost scales as $O(2^c)$ rather than with the total number of full path combinations. The practical efficiency of IPA therefore depends primarily on the number and nesting depth of such dependency-inducing structures. The next section introduces a progressive decomposition strategy that reduces repeated conditioning by replacing resolved diamond subgraphs with statistically equivalent supernodes.

4. Progressive Decomposition of Diamond Structures

Progressive decomposition is the mechanism that makes IPA practical on many non-trivial DAGs with reconvergent structure. It preserves exactness by replacing resolved diamonds with conditionally equivalent supernodes,

Algorithm 1 Information Propagation Algorithm (IPA)

```

1: Compute a topological ordering of  $G$ 
2: for each node  $v$  in topological order do
3:   if  $v \in S$  then  $b(v) \leftarrow R(v)$ 
4:   else
5:     Compute  $S_i = b(p_i) P(p_i, v)$  for all  $p_i \in \text{Pa}(v)$ 
6:     if incoming parent signals are independent then
7:        $b(v) \leftarrow R(v)(1 - \prod_i (1 - S_i))$ 
8:     else
9:       Identify the reconvergent (diamond) structure affecting  $v$ 
10:      Resolve its dependency by conditional enumeration on the separating set
11:      Compute  $b(v)$  from the conditioned signal probabilities
12:    end if
13:  end if
14: end for
15: return  $\{b(v)\}_{v \in V}$ 

```

and it improves efficiency by preventing the same dependencies from being conditioned on repeatedly.

While conditional decomposition resolves dependencies locally, real process networks rarely contain only isolated diamonds. In practical directed acyclic graphs, reconvergent structures often overlap or nest within one another, so that the output of one diamond becomes part of a larger downstream reconvergent structure. In such cases, repeatedly conditioning on the same upstream variables can lead to severe computational redundancy. IPA addresses this issue through progressive diamond decomposition. Once a diamond subgraph has been solved, it is replaced by a statistically equivalent supernode that preserves its exact conditional behaviour while simplifying the remaining graph.

More generally, for a network with L nested dependency levels, a direct conditional evaluation may require

$$b(j) = R(j) \sum_{s_1 \in \{0,1\}^{m_1}} \sum_{s_2 \in \{0,1\}^{m_2}} \cdots \sum_{s_L \in \{0,1\}^{m_L}} P(s_1) P(s_2 | s_1) \cdots P(s_L | s_1, \dots, s_{L-1}) \phi(s_1, \dots, s_L), \quad (13)$$

where m_ℓ is the number of conditioning nodes introduced at nesting level

ℓ , and $\phi(\cdot)$ is the resulting signal probability at the join node. This cost grows exponentially with both the number of conditioning variables and the depth of nesting. The purpose of progressive decomposition is to avoid this repeated re-evaluation by solving inner diamonds once, storing their exact conditional behaviour, and reusing that information whenever the reduced network is evaluated further downstream.

4.1. Supernode Representation and Exactness Guarantees

Let D be a diamond subgraph with fork-side conditioning set $C_D = \{c_1, \dots, c_m\}$ and join node j_D . Solving the diamond yields the exact conditional signal probability at the join node for every configuration of the conditioning variables:

$$\psi_D(s) = P(J_D \mid s), \quad s \in \{0, 1\}^m, \quad (14)$$

where J_D denotes the event that a signal reaches the join node of the diamond and s is a binary state assignment to C_D . The resolved diamond is then replaced by a supernode that stores the map $s \mapsto \psi_D(s)$. This supernode acts as a statistically equivalent component in the reduced graph: whenever the conditioning state s is encountered in a larger downstream computation, the contribution of the original diamond is recovered exactly through (14), without re-solving the full subgraph. The following lemmas make the exactness guarantees precise.

Lemma 1 (Conditional invariance). *Let $A \subseteq V$ and let $\mathcal{R}_A$ denote the vector of reachability indicators of A . Then for any node v ,*

$$b(v) = \sum_{c \in \{0,1\}^{|A|}} P(\mathcal{R}_A = c) P(\mathcal{R}_v = 1 \mid \mathcal{R}_A = c), \quad (15)$$

and the decomposition is well defined because every $\mathcal{R}_u$ is a deterministic function of the mutually independent component indicators $\{X_w, X_{wu}\}$. Moreover, let D be a diamond subgraph with join j and conditioning set $C \subseteq F \cap \text{anc}(j)$, and let $\psi_D(c) = P(J_D \mid \mathcal{R}_C = c)$ denote its conditional signal map. If an enclosing computation conditions on further variables A' with $A' \cap D = \emptyset$ — that is, the outer conditioning fixes no variable internal to D — then $\psi_D(c)$ is unchanged: outer conditioning affects only the weights $P(\mathcal{R}_C = c \mid \mathcal{R}_{A'})$, never the map ψ_D itself.

Proof. The first display is the law of total probability over the partition $\{\mathcal{R}_A = c\}$. For the second claim, fix c . Conditional on $\mathcal{R}_C = c$, the event J_D is determined by the indicators of nodes and edges internal to D . By Assumption 2 these are independent of every indicator outside D , hence of $\mathcal{R}_{A'}$ whenever A' contains no internal variable of D . Consequently $P(J_D \mid \mathcal{R}_C = c, \mathcal{R}_{A'} = a) = P(J_D \mid \mathcal{R}_C = c)$ for every outer state a . $\square$

The condition $A' \cap D = \emptyset$ is essential and is enforced by the identification procedure: when an outer conditioning does fix a variable internal to a previously resolved subgraph — as can occur when diamonds overlap or share conditioning nodes — the sub-problem is a different one and is resolved (and stored) separately under its own conditioning context. Section 4.2.2 makes this identification precise. A conditioning variable already fixed by an outer state contributes a single term to the enumeration rather than two, so overlapping conditioning never doubles work unnecessarily.

Lemma 2 (Separator sufficiency). *Let v be a join node with parents $\text{Pa}(v)$, and let $C \subseteq F \cap \text{anc}(v)$ be such that every directed path from a common ancestor of two distinct parents $u_i, u_j \in \text{Pa}(v)$ to those parents intersects C — that is, C is a cutset of the parents' shared ancestry. Then, conditional on $\mathcal{R}_C$, the parent reachability events $\{\mathcal{R}_{u_i}\}$ are mutually independent, and*

$$b(v \mid \mathcal{R}_C = c) = R(v) \left(1 - \prod_{u \in \text{Pa}(v)} (1 - P(\mathcal{R}_u = 1 \mid \mathcal{R}_C = c) P(u, v)) \right). \quad (16)$$

Proof. With $\mathcal{R}_C$ fixed, the residual ancestries $\text{infl}(u_i; C)$ of distinct parents share no component indicator, by the cutset property. Each $\mathcal{R}_{u_i}$, conditional on $\mathcal{R}_C$, is a function of its residual ancestry alone; functions of disjoint sets of independent variables are independent. The display is then the independent-union rule (8) applied conditionally. $\square$

Lemma 2 is the justification for the conditioning sets used throughout: the identification procedure returns, for each reconvergent structure, a separating set C satisfying the cutset property, and the per-structure cost $2^{|C|}$ follows.

Lemma 3 (Independent-substructure factorisation). *Fix a conditioning context E_c and, for parents u, u' of a join v , write $u \sim u'$ whenever $\text{infl}(u; E_c) \cap \text{infl}(u'; E_c) \neq \emptyset$. Let $G_1, \dots, G_m$ be the equivalence classes of $\text{Pa}(v)$ under the*

transitive closure of $\sim$. Then the events $\{v \text{ is reached through some parent in } G_k\}$, $k = 1, \dots, m$, are mutually independent, so

$$b(v) = R(v) \left(1 - \prod_{k=1}^m (1 - P(v \text{ reached via } G_k)) \right), \quad (17)$$

and each class G_k can be conditioned on its own separating set C_k independently of the others.

Proof. Distinct classes have disjoint unresolved influence by construction, so the class events are functions of disjoint sets of independent indicators, hence independent; the product form is the union rule for independent events. $\square$

The practical effect is substantial. At a join whose k parents share ancestry only pairwise within m independent groups, joint conditioning would enumerate $2^{|C_1| + \dots + |C_m|}$ states, whereas factorised conditioning enumerates $\sum_k 2^{|C_k|}$. In the extreme case of k parents each depending on its own single fork, the cost falls from 2^k to $2k$ — from exponential to linear in the fan-in.

Lemma 4 (Supernode equivalence). *Let D be a resolved diamond with conditional signal map $\psi_D(s)$. Replacing D by a supernode storing $\psi_D(s)$ preserves the exact source-to-node reliability of every downstream node in the reduced graph.*

Proof. For any downstream computation, the contribution of D enters only through the probability that its join node receives a signal under a given conditioning state. By Lemma 1 the stored map $\psi_D(c)$ equals $P(J_D \mid \mathcal{R}_C = c)$ in every outer context that fixes no internal variable of D , and by Lemma 2 the downstream update uses only these conditional probabilities together with the weights $P(\mathcal{R}_C = c)$; both are preserved by the replacement, so every downstream belief is unchanged. $\square$

4.2. Progressive Decomposition Procedure

IPA applies diamond decomposition progressively during topological evaluation. When the algorithm reaches a join node involved in a reconvergent structure, it: identifies the associated diamond subgraph; determines the conditioning variables required to resolve its internal dependencies; computes the conditional signal map of the join node; and replaces the entire resolved diamond by a supernode in the reduced graph. This procedure transforms the

Table 1: Trace of the worked multi-level example (Figure 5): the four diamond resolutions in their base context. $w = P(\mathcal{R}_f = 1)$; $\psi(1), \psi(0)$ are the conditional signal probabilities at the join given the fork reachable/unreachable.

Diamond	Join j	Fork f	w	$\psi(1)$	$\psi(0)$	$b(j) = w \psi(1) + (1-w) \psi(0)$
D_1	6	1	0.81	0.85910	0.53144	0.79684
D_2	11	7	0.81	0.80437	0.43047	0.73333
D_3	12	14	0.81	0.83867	0.52496	0.77907
D_4	16	1	0.81	0.82255	0.73623	0.80615

original DAG into a progressively simpler equivalent DAG. Inner diamonds are resolved first, so that larger outer diamonds can treat them as supernodes rather than re-evaluating their internal paths. Figure 5 illustrates the decomposition of a nested reconvergent network.

The topological order ensures that the inner diamond D_1 is detected and resolved before downstream join nodes are processed. Once replaced by a supernode, the reduced graph contains fewer explicit dependency-inducing paths. The same procedure is then applied to D_2 and D_3 , after which the final outer diamond D_4 can be evaluated using the already-resolved supernodes. In this way, dependencies are handled locally and only once.

4.2.1. Worked multi-level example

To make the procedure concrete, consider the network of Figure 5 with all non-source nodes assigned intrinsic reliability 0.9, all edges transmission probability 0.9, and source reliabilities 1.0. Table 1 traces the four diamond resolutions: for each, the fork whose reachability is conditioned on, the weight $w = P(\mathcal{R}_f = 1)$, the two conditional signal maps, and the resulting belief at the join node. The inner diamond D_1 (fork 1, join 6) is resolved by enumerating the two states of its fork; note that its join retains a nonzero conditional value when the fork is down, because the path from source S_2 through node 13 is unaffected by the fork. D_2 and D_3 are resolved in the same way. The outer diamond D_4 (fork 1, join 16) then conditions on the same fork; within each of its two states, the nested structures at joins 6, 11, and 12 are required under that outer context and are supplied by their stored conditional maps — re-resolved only where the outer state changes their context (Section 4.2.2), for example $b(12 \mid \mathcal{R}_1=1) = 0.8098$ versus $b(12 \mid \mathcal{R}_1=0) = 0.6481$.

The final belief at node 16 is $b(16) = 0.80614907$, and every node’s belief




Figure 5: Progressive diamond decomposition in a nested reconvergent network.

agrees with exhaustive path enumeration (inclusion–exclusion over all source-to-node paths) to within 3.4×10^{-16} , i.e. floating-point round-off. The effect of reuse is also measurable on this small example: resolving the network requires the evaluation of 68 distinct conditional sub-problems with supernode reuse, against 145 evaluations if every nested structure is re-evaluated in every context in which it appears.

4.2.2. Identification of stored sub-problems

A resolved sub-structure is identified by the pair (induced substructure, conditioning context): the internal nodes and edges of the diamond together with the set of upstream variables whose states are fixed when it is resolved. Two occurrences of the same substructure reached under different conditioning contexts are treated as distinct sub-problems, because an outer state that fixes a variable inside the substructure changes its conditional signal map (cf. the condition of Lemma 1). This context-aware identification is what guarantees correctness when diamonds overlap or share conditioning variables: reuse occurs only between genuinely identical sub-problems. Conversely, a conditioning variable whose state is already determined by the enclosing enumeration contributes a single state rather than two, so the enumeration never expands over variables that carry no remaining uncertainty. Together these two rules bound the stored state space by the number of genuinely distinct conditional sub-problems, which is the quantity reported as realised computational work in Section 5.2.

4.3. Complexity and Relation to Established Exact Methods

The computational cost of the procedure admits an exact per-instance expression. Let the recursion of Section 4.2 resolve diamonds $d = 1, \dots, M$ (counting nested resolutions), with conditioning sets C_d and internal edge sets E_d . The work performed is

$$W = \sum_{d=1}^M 2^{|C_d|} \cdot O(|E_d|), \quad (18)$$

since each diamond is resolved by enumerating the $2^{|C_d|}$ states of its conditioning set and each state requires a propagation over the diamond’s internal structure. Every term of this sum is available before enumeration begins, because the conditioning sets are produced by the identification procedure itself; the cost of an analysis is therefore computable in advance for a given

network, rather than merely bounded asymptotically. Memoisation of resolved sub-problems (Section 4.2.2) reduces the realised work to the number of genuinely distinct conditional sub-problems, which is the quantity measured in the experiments below. The factorisation of Lemma 3 reduces the exponents themselves in favourable cases — at a fan-in of k independent branches, from 2^k states to a number linear in k — without changing the worst-case class.

Relation to cutset conditioning and junction trees. The procedure is a specialisation of cutset conditioning (Pearl 1988) to source-to-node reachability in DAGs. Classical cutset conditioning selects a single global cutset whose state enumeration renders the whole network singly connected, at cost exponential in the cutset size; the present method instead identifies, at each reconvergence actually encountered, a local separating set of the parents’ shared ancestry (Lemma 2), conditions on it, and recurses, so that the realised cost is a sum of local exponentials rather than one global exponential. Tree-structured regions of the network are traversed by the closed-form updates of Section 3 at linear cost, and the factorisation of Lemma 3 prevents joint conditioning over forks that are independent given the current context. The maximum conditioning-set size encountered on a network is bounded by the size of the largest set of mutually entangled shared ancestors, a quantity that is in turn bounded by — and, as shown below, empirically tracks — the treewidth of the network. The method therefore sits in the same parametric complexity class as the junction-tree algorithm (cost exponential in treewidth; Lauritzen and Spiegelhalter 1988) and as a well-ordered reduced binary decision diagram (size exponential in pathwidth): all are exact, all are fixed-parameter tractable in the relevant width, and none escapes the #P-hardness of the underlying problem. No closed-form complexity in the node count alone exists for any of them.

Measured comparison. Section 5.2 reports, for every network in the validation corpus, the maximum conditioning-set size, the realised number of conditional sub-problems, and the node count of a reduced-ordered binary decision diagram for the same network built under dynamic variable reordering. Two findings anticipate the details. First, the maximum conditioning-set size sits in the same band as the logarithm of the decision-diagram size — for example 7 against 8.2 on a 4×4 grid, 10 against 9.0 on a bridge network, 13 against 11.7 on the densest random network tested (Figure 7) — confirming

that both methods are governed by the same structural width parameter. Second, the realised work of the proposed method is of the same order as the decision-diagram size, sometimes smaller and sometimes larger, with neither method dominating across the corpus. We state the conclusion plainly: for exact point-valued reliability the proposed method is competitive with, but not structurally superior to, a well-ordered decision diagram. Its distinguishing capability lies in the native propagation of imprecise component reliabilities (Section 5.3), which decision-diagram methods do not provide.

4.4. Numerical Realisation

The propagation scheme of Sections 3–4 is arithmetic over the operations $\{+, \times, \text{complement}\}$ on component reliabilities, and is realised once, generically, for the three input forms of Section 2.5: point values, intervals, and probability boxes. Three numerical points warrant note. First, multi-parent updates are evaluated in the product form $1 - \prod_i (1 - p_i)$ rather than by the alternating inclusion–exclusion sum; the two are equivalent under independence, but the product form is numerically stable and, for interval-valued inputs, avoids the spurious widening that alternating sums introduce. Second, for interval inputs the conditioning step is evaluated at the corner states of the conditioning bounds and the result enclosed in an interval envelope; Section 5.3.1 shows this yields the exact range, by monotonicity. Third, for p-box inputs the conditioning recombination at a join is a convex combination of two dependent random quantities, not an independent sum; it is evaluated by the operator of Section 5.3.2, which integrates over the distribution of the conditioning weight and bounds the branch dependence by the Fréchet–Hoeffding inequalities. Because the underlying floating-point arithmetic is not outward-rounded, the computed bounds are intersected with $[0, 1]$ as a final step — a projection that can only tighten a bound on a probability, never invalidate it — and this projection is guarded: an excursion beyond $[0, 1]$ exceeding ordinary floating-point round-off by several orders of magnitude raises an error rather than being silently discarded, so that a genuine soundness regression cannot be masked as rounding. The full implementation, together with all validation scripts and data, is openly available (see Data and Software Availability).

5. Case Studies

This section evaluates IPA in four steps: a benchmark validation on a grid network (Section 5.1); a corpus-wide validation and cost comparison against an independent exact decision-diagram method (Section 5.2); the imprecise-propagation capability (Section 5.3); and an applied reliability case study of a medical drone logistics network for Scotland (Section 5.4). All runtimes reported in this section were measured after a warm-up run of the same computation was executed and discarded, so that one-off program-initialisation cost is excluded from every figure; each reported time is the minimum of repeated measurements on otherwise idle hardware.

5.1. Benchmark Validation on a Grid Network

5.1.1. Benchmark setup

Both the proposed Information Propagation Algorithm and Tong and Tien’s directed Probability Propagation Method (dPrPm) (Tong and Tien 2019) compute node reliability in directed acyclic networks by propagating local probability information from upstream to downstream nodes. Both methods exploit the acyclic structure of the graph to ensure that upstream dependencies are processed before downstream updates. The two methods differ in how dependencies are represented and resolved. dPrPm propagates both marginal and pairwise probabilities and uses local approximations to account for dependent interactions. Specifically, when multiple predecessors influence a node, the update rules explicitly rely on approximating a three-node joint distribution for every new edge introduced. Although effective, this approach requires storing both marginal and pairwise probabilities throughout the network, which can become memory-intensive in large or dense networks. IPA takes a fundamentally different view of network dependencies: it identifies reconvergent dependency patterns as diamond structures and resolves them through conditional decomposition. Rather than tracking pairwise relationships, it decomposes these structures into statistically equivalent supernodes, so the relevant dependency structure is captured inherently in the decomposition.

To provide a direct benchmark, IPA is applied to the same 16-node directed grid network studied in Tong and Tien (2019), shown in Figure 6. This benchmark is useful because it contains multiple reconvergent structures while remaining small enough for exact verification by explicit path-based computation. Following the source study, all source nodes are assigned




Figure 6: 16-node directed grid network (Tong and Tien 2019). Source nodes 1, 3, and 13 are coloured in grey.

prior reliability 1.0 and all directed edges are assigned transmission probability 0.9.

5.1.2. Accuracy results

Table 2 compares the node reliabilities obtained using IPA with exact values and with the lower and upper bounds reported for dPrPm in Tong and Tien (2019). IPA reproduces the exact values up to floating-point precision. By contrast, dPrPm provides tight lower and upper bounds that coincide with the exact values for many nodes in this benchmark.

Table 3 reports the runtime of IPA on the benchmark network alongside the runtime reported for dPrPm in Tong and Tien (2019). Two caveats attend this comparison and we state them plainly. First, the dPrPm figure is quoted from the source study and was obtained in a different implementation environment; the comparison is indicative only, and no controlled speed claim is based on it. Second, the dPrPm reference implementation is not publicly available and its authors could not be reached, so the comparison cannot be reproduced independently. For these reasons the quantitative performance assessment of the proposed method in this paper rests on the controlled,

Table 2: Node reliabilities on the benchmark grid network. IPA reproduces the exact values; dPrPm bounds are as reported in Tong and Tien (2019). *[Author note: carry the dPrPm bound columns over from the original submission’s table.]*

Node	Exact = IPA	Node	Exact = IPA
1	1.0	9	0.981
2	0.99	10	0.8829
3	1.0	11	0.97734390
4	0.98015430	12	0.97456525
5	0.9	13	1.0
6	0.99510274	14	0.979461
7	0.98955925	15	0.98497989
8	0.89060332	16	0.98538853

Table 3: Computational performance on the benchmark grid network. The dPrPm runtime is quoted from Tong and Tien (2019) and is indicative only (different implementation environment; memory data not available).

Method	Runtime	Memory
IPA	2.378 ms	4.94 MB
dPrPm (reported)	0.32 s	n/a

same-environment comparison against an independent exact method in Section 5.2, and dPrPm is retained only as the published point of reference that motivated the benchmark.

The benchmark has also been re-verified independently. The node reliabilities produced by the proposed method were compared against a reduced-ordered binary decision diagram with dynamic variable reordering — a canonical independent exact method — constructed for the same network; the worst per-node disagreement is 1.1×10^{-16} , i.e. floating-point round-off, confirming the values in Table 2. During this hardening the diamond-identification procedure was also stress-tested beyond the benchmark’s parameterisation (non-unit node priors, and the boundary case in which all node priors equal one); two edge-case identification errors were found and corrected, and the corrected procedure is the one validated across the full corpus of Section 5.2. Neither error affects the benchmark table, which stands as originally reported.

Table 4: Exactness and complexity of IPA on named structured networks, versus the exact ROBDD oracle. *root D./uniq. D.* are the numbers of root and unique diamond structures identified; $\max|C|$ is the largest conditioning set; *ROBDD* is the node count under dynamic reordering; $\max|\Delta|$ is the worst per-node deviation of IPA from the exact ROBDD reliability.

Network	$ V $	$ E $	src	root D.	uniq. D.	$\max C $	ROBDD	$\max \Delta $
power network	23	27	3	5	9	3	384	8.3×10^{-17}
grid (benchmark)	16	24	3	8	41	7	290	1.1×10^{-16}
Karl network	26	74	1	18	147	11	1921	5.6×10^{-17}
counterexample-n15	15	23	1	4	13	4	340	5.6×10^{-17}

5.2. Validation and Cost Comparison Against an Exact Decision-Diagram Method

To address the breadth of validation, the proposed method was compared against an independent exact method — a reduced-ordered binary decision diagram (ROBDD) with dynamic variable reordering — across a corpus of 129 random and mutated DAGs (10 to 28 nodes, densities 0.14–0.42), six topological families (multi-source, grid/lattice, layered, bridge, series-parallel, complete), larger random networks up to 50 nodes, and several real infrastructure networks. Every network in the corpus was evaluated by both methods; the worst per-node disagreement observed anywhere is 1.1×10^{-16} , for both perfect and imperfect component reliabilities. Table 5 summarises the random corpus and Table 4 reports representative structured and real networks in detail.

The comparison also quantifies cost. For each network we report the maximum conditioning-set size $|C|_{\max}$ encountered by the proposed method, its realised number of distinct conditional sub-problems, and the ROBDD node count under dynamic reordering. As anticipated in Section 4.3, $|C|_{\max}$ tracks the logarithm of the ROBDD size across the corpus (Figure 7) — the two methods are governed by the same structural width — and the realised work of the two methods is of the same order, with neither dominating: on some networks the proposed method resolves fewer sub-problems than the diagram has nodes (e.g. 42 against 340 on the counterexample network; 142 against 759 on the complete 8-node network), on others more (e.g. 1961 against 1043 on a layered network). No family in the corpus favours the proposed method for exact point-valued computation. This is an honest structural finding, and it sharpens rather than weakens the contribution: the

Table 5: Random and mutated DAG corpus (129 unique graphs), summarised by family. Random DAGs are generated per family with edge probability $p \in \{0.10, 0.15, 0.20\}$ and eight seeds; *mutant-n28* is a base $n=28$ DAG with $+5/-2$ edge mutations. Every graph in every family was validated exact against the ROBDD oracle ($\max |\Delta| \leq 1.1 \times 10^{-16}$ across the whole corpus). *Naive blow-ups* are builds exceeding 2.5×10^6 ROBDD nodes under a naive (topological) variable order, all of which dynamic reordering resolves.

Family	#graphs	$ V $	$ E $	density	med. ROBDD	max. ROBDD	naive blow-ups
random-n10	24	10	12–18	0.27–0.40	80	270	0
random-n12	24	12	15–28	0.23–0.42	197	693	0
random-n15	24	15	21–35	0.20–0.33	450	1092	0
random-n20	24	20	31–50	0.16–0.26	1164	2849	0
random-n25	24	25	42–73	0.14–0.24	3119	20323	9
mutant-n28	8	28	62	0.16	2562	6050	6
counterexample	1	15	23	0.22	340	340	0

value of the proposed method rests on the capabilities of Section 5.3, not on a claim of a smaller exact state space.

5.3. Imprecise Reliability Propagation

The results so far establish that the proposed method matches a well-ordered decision diagram on exact point-valued reliability at comparable cost. This section presents the capability that separates the two: the propagation scheme accepts interval- and p-box-valued component reliabilities directly (Section 2.5).

5.3.1. Interval-valued reliabilities: exact and fast

Because the node reliability $b(v)$ is non-decreasing in every component reliability (Section 2.3), its exact range over interval inputs is attained at the two extreme corners of the input box: the lower bound when every component takes its lower reliability, the upper bound when every component takes its upper reliability. No dependence issue arises, and the propagated interval is exact. This was verified across the full corpus of Section 5.2 by comparing the propagated interval against the decision-diagram evaluation at the two corner inputs: the worst discrepancy in either bound is of order 10^{-16} . By contrast, interval propagation that ignores reconvergent dependence over-widens the range — by up to 0.45 in the corpus — illustrating that the conditioning machinery is as necessary for sound imprecise propagation as it is for exact point values.




Figure 7: The proposed method’s maximum conditioning-set size against the logarithm of the sifted ROBDD node count, across the complexity-validation networks. Both quantities measure the same structural width of the network; neither method dominates.

The interval computation is also fast. For a one-shot query — given a network, produce the interval reliability — the proposed method completed faster than the decision-diagram route to the same answer (build the diagram once, evaluate at both corner inputs) on every one of eight corpus families tested, by factors ranging from 2.9 to 95 (Table 6). We scope this claim carefully: it concerns the one-shot cost, in which the diagram’s construction cost counts against it. A workload that amortises a single construction across many repeated queries on a fixed network is a different scenario, not tested here, and could favour the diagram-based route instead.

5.3.2. Probability-box reliabilities: guaranteed distributional bounds

For p-box inputs the conditioning step requires care. At a join resolved by conditioning, the belief is a convex combination $W \cdot A + (1 - W) \cdot B$, where W is the (uncertain) probability of the conditioning state and A, B are the (uncertain) conditional beliefs of the two branches. A and B are functions of shared upstream reliabilities and are therefore dependent; treating the combination as an independent sum is unsound and can assign probability mass outside $[0, 1]$. The operator used here instead integrates over the distribution of W , blending the branch distributions at each weight level, with the branch

Table 6: One-shot interval reliability: a single propagation of the proposed method against the decision-diagram route (build once, evaluate at both corner inputs), warm-measured. The proposed method is faster on every family tested; the comparison is scoped to one-shot queries.

Network	$ V $	$ E $	ROBDD	IPA (ms)	ROBDD route (ms)	ratio
counterexample	15	23	323	0.81	52.3	64.5
multi-source	25	57	5622	22.0	193.8	8.8
grid 4×4	16	21	248	0.58	55.0	95.1
layered 4×6	23	35	484	4.00	58.4	14.6
bridge	16	25	529	13.3	48.3	3.6
series-parallel	23	32	215	2.66	64.1	24.2
complete-8	8	28	685	1.58	49.4	31.2
random $n=25$	25	53	4146	40.3	117.2	2.9

dependence bounded by the Fréchet–Hoeffding inequalities (Williamson and Downs 1990). Soundness of the resulting bounds is inherited from that classical result: the computed pair of distribution functions encloses the true distribution of the node reliability under any dependence between the branches. A tighter variant restricts the dependence bound to non-negative dependence, motivated by the observation that both branches are non-decreasing functions of the shared upstream reliabilities and hence cannot be negatively dependent; this variant produced no soundness violation in any of the validation configurations tested (dozens of topology, input-shape, and uncertainty-width combinations, each checked against Monte Carlo ground truth), but a general proof of its soundness under partial dependence information of this kind is not available in the literature and is identified as future work. All headline claims in this paper use only the provable Fréchet-based bound.

The tightness of the bounds — as distinct from their soundness, which does not vary — depends on the structure: across a sweep of input-distribution shapes, uncertainty widths, and reconvergence regimes on the benchmark network, the bound band ranged from approximately 0.18 of the unit range under high-reliability, weakly reconvergent conditions to approximately 0.70 under strongly reconvergent, fully uncertain conditions, with zero soundness violations across the sweep (Figure 8). The driver of tightness is reconvergence depth combined with input uncertainty; the shape of the input distribution is secondary. Computational cost is a controllable dial: interval propagation costs a small constant factor over point-valued propagation (a




Figure 8: Tightness envelope of the certified p-box bound on the benchmark network: band width against component-reliability uncertainty width, for three input-distribution shapes and two reconvergence regimes. Soundness holds in every configuration (zero violations); only tightness varies, driven by the regime rather than the input shape.

factor of about 1.2 in our measurements: 0.67 ms against 0.79 ms on the 15-node reference network), while p-box propagation grows quadratically with the discretisation level of the distribution bounds (2.7 s, 8.3 s, and 110 s at 50, 200, and 800 levels respectively on the same network), in exchange for proportionally tighter bands. We therefore state bound tightness always at the discretisation level used, and make no claim that distributional propagation is cheap — it is controllable, with rigorous bounds at every setting.

5.3.3. A certified decision bound from a single propagation

The practical payoff is that the propagated distribution bounds certify decision-relevant probabilities. Consider a stated reliability requirement x^* for a given node: the propagated p-box yields guaranteed lower and upper bounds on $P(b(v) \leq x^*)$, the probability that the node fails its requirement, from one propagation. On the benchmark network at a representative requirement of $x^* = 0.95$, the certified band on this probability had width between 0.02 and 0.12 across the regimes tested (Table 7). Matching that precision by simulation, planned conservatively without foreknowledge of the answer, would require between 267 and 9,604 samples of the network’s un-

Table 7: Certified bound on the requirement-violation probability $P(b(v) \leq 0.95)$ at the benchmark network’s target node, from a single propagation, against the Monte Carlo sample count required to match the certified precision under worst-case a-priori planning.

Regime	Uncertainty width	Certified band	Band width	MC samples (planning)
perfect components	0.05	[0.000, 0.020]	0.02	9,604
perfect components	0.10	[0.000, 0.100]	0.10	385
perfect components	0.15	[0.000, 0.120]	0.12	267
uncertain components	0.05–0.15	[0.980, 1.000]	0.02	9,604

certain inputs, and the result would remain a statistical estimate rather than a guarantee. (When the true probability lies near 0 or 1 the retrospective sample requirement collapses — correct statistics, but knowable only in hindsight; the planning figure quoted is the one a practitioner must budget.) An exact decision-diagram method has no analytic route to this bound at all and would itself fall back on simulation. This certified-bound capability, not raw speed, is the concrete engineering value of imprecise propagation.

5.4. Applied Case Study: A Medical Drone Logistics Network for Scotland

The methodology having been validated in Sections 5.1–5.3, this section applies it to a real infrastructure planning problem: the conceptual medical drone delivery network for Scotland designed by Jones et al. (2025), who optimised station placement and network structure over real hospital, airport, and candidate-station locations, two drone types (short-range vertical take-off and landing, nominal range 70 km; long-range fixed-wing, nominal range 700 km), and stated infrastructure specifications, trading off delivery time, capital cost, and resilience. The reliability question the present method can answer, and their point-probability framework cannot, is: given the acknowledged uncertainty in the component availability assumptions, how confidently can each location in a candidate design be said to be reachable from the supply hubs?

5.4.1. Model inputs and their provenance

Every input of the reliability model is derived from the source study, with one clearly flagged extension.

Location availability. The source study’s own resilience analysis assumes that non-hub locations fail independently with probability 0.2 while hub locations

are always available. We adopt the same structure, but represent the non-hub availability as the interval $[0.75, 0.85]$ centred on the source study’s figure, for a deliberate reason: that figure is stated in the source study as an assumption rather than derived from data, and an interval-valued input turns the analysis into a built-in sensitivity assessment of precisely that assumption. Hub locations retain availability 1.

Connection existence. A connection between two locations exists, for a given drone type, if and only if the inter-location distance is within that type’s stated nominal range — the source study’s own network-construction rule, applied exactly. Where both drone types can serve a connection, the connection is available if either type’s transmission succeeds, the two being independent.

Connection reliability — the flagged extension. The source study holds weather conditions constant in its own analysis, while noting that weather-driven range variation is in principle an uncertain quantity requiring expert elicitation, which it deliberately defers. We extend exactly this deferred point, and only this point. A connection operating well within a drone’s nominal range is treated as reliable. A connection whose length falls within an illustrative weather-derating margin of the range limit — the final ten per cent, an assumed bound since the source study does not quantify one either — is treated as genuinely uncertain: its transmission probability is the vacuous interval $[0, 1]$, expressing honestly that adverse weather may or may not close that particular route. This is an extension of the source study’s own stated limitation, not an independent invention, and no other probability in the model is assumed.

Direction of dependency. The source study’s connections are undirected, since a delivery route can be flown in either direction; the reliability question, however, concerns supply reaching facilities from hubs. The dependency graph is therefore directed from the hub tier outward — matching the source study’s own hub-and-spoke architecture — and source-to-node reachability is computed on the resulting acyclic structure.

5.4.2. Network configurations

The source study’s actual optimised network layouts are not published. Three configurations are therefore constructed as proxies for the qualitative character of three trade-off points described in its results: a fixed-wing-reliant

centralised design serving remote locations through a hub backbone (217 locations, 263 connections); a densely interconnected short-range design (242 locations, 1,753 connections); and a minimal-investment design retaining only pre-existing, non-optional infrastructure (230 locations, 1,648 connections). Each is explicitly a proxy for the described character of the corresponding design, not a reproduction of an unavailable optimisation result.

5.4.3. Redundancy and the practical boundary of exact analysis

A network-design parameter controls the reconvergent structure of each configuration: the number of alternate routes provisioned per location — a genuine redundancy choice in network design, not a parameter of the analysis method. Connecting every operationally plausible pair of locations produced conditioning-set requirements of 27–28, well beyond the practical range of exact analysis (the original submission reported computation becoming impractical beyond a nesting depth of approximately 18, a figure the reviewers also highlighted). Bounding the provision to sixteen alternate routes per location brought the requirement to 15–17 — just inside that range — and further increases in provisioned routes changed neither the requirement nor the runtime materially, indicating that the reconvergent structure of this network saturates at that level. The full exact interval propagation of the densest configurations completes in under 25 seconds; the sparse centralised configuration, whose conditioning requirement is 10, completes in well under one second. The practical boundary of exact analysis on this real network is thus a measured quantity, expressed in terms of a controllable design parameter: more provisioned redundancy improves resilience but raises the cost of verifying it exactly.

The comparison of Section 5.2 was repeated on these networks. On the sparse centralised configuration the two methods agree to floating-point precision, and the one-shot interval computation of the proposed method is approximately fourteen times faster than the decision-diagram route to the same interval (0.012s against 0.173s; diagram size 2,667 nodes). On the minimal-investment topology provisioned at low redundancy (six alternate routes per location), both methods are comfortably fast and the proposed method remains moderately faster (approximately six times: 0.128s against 0.830s; diagram size 4,968 nodes; agreement 1.1×10^{-16}). On the same topology at sixteen alternate routes per location, the decision-diagram build did not complete within a practical computational budget under either of two variable-ordering strategies (dynamic reordering, and a fixed topologi-

cal order), while the proposed method completed the full exact computation in under 25 seconds. We state the finding precisely: both methods are governed by the same underlying network property, and both are practical at low redundancy; on this network, the practical range of the proposed method, measured along a real redundancy design parameter, extends further than that of the decision-diagram alternative. No claim is made that decision-diagram methods cannot handle this problem class in general, nor that the crossover point has been located precisely.

5.4.4. *Reliability findings*

Table 8 summarises the propagated interval reliabilities. They are informative, not vacuous, and they differ meaningfully across the three designs. In the centralised, low-redundancy configuration the propagated bounds are narrow (band width at most 0.10, mean 0.089): each location’s reachability essentially inherits the availability band assumed for the locations themselves, with little additional uncertainty accumulated through the network — but also little protection: the design is tree-like, so single-route failures directly disconnect service. In the two higher-redundancy configurations the best-connected locations are reachable with near-certainty, while the worst-served locations carry a genuine, decision-relevant spread: the least certain facility — Islay Hospital, an island site — has reachability probability bounded between 0.56 and 0.72. That interval is the honest answer to the planning question under the acknowledged input uncertainty: no point-probability analysis of the same design could distinguish a facility whose reachability is confidently 0.6 from one whose reachability is anywhere between 0.56 and 0.72, yet the two call for different interventions. Figure 9 maps the network coloured by the lower reliability bound and by band width, identifying the specific facilities — concentrated in the islands and the far north — where additional redundancy, or better component-availability data, would most improve either the network or our knowledge of it.

5.5. *Limitations and Future Work*

The method is exact, and we are explicit about where its advantages do and do not lie.

Exact point-valued reliability. The method offers no asymptotic advantage over a well-ordered decision diagram. Both are exponential in the network’s

Table 8: The three proxy configurations of the applied case study: structure, conditioning requirement, full interval-propagation soundness, and the propagated reliability bounds. Every belief is contained in $[0, 1]$ with $\underline{b} \leq \bar{b}$ (soundness); the worst-served facility in both dense configurations is Islay Hospital.

Configuration	$ V $	$ E $	$ C _{\max}$	mean band	max band	worst facility
fixed-wing centralised	217	263	10	0.089	0.100	$[0.750, 0.850]$
dense decentralised	242	1753	17	0.094	0.161	$[0.562, 0.722]$
concentrated minimal	230	1648	16	0.093	0.161	$[0.562, 0.722]$

structural width — the conditioning width here, the pathwidth for the diagram — reflecting the #P-hardness of network reliability; across the corpus neither dominates, and the realised work tracks the diagram size to within a constant factor (Section 5.2). The method is competitive with, not superior to, a well-ordered decision diagram for exact point evaluation; its distinctive value lies in imprecise propagation.

Practical range. Being exact, the method is exponential in width and intended for the sparse-to-moderate regime — a ceiling shared with every exact method. The applied case study makes the boundary concrete on one real network: conditioning requirements of 15–17, reached at sixteen provisioned alternate routes per location, remain practical (tens of seconds); the requirement of 27–28 produced by unrestricted connectivity does not. Imprecise propagation does not relax this boundary: interval and p-box propagation use the same conditioning depth as exact point propagation, and only a change to the network design itself alters the cost. For networks beyond the practical range, a depth-limited hybrid suggests itself: condition exactly to a chosen depth and, beyond it, combine reconvergent contributions without conditioning, which for interval inputs yields a rigorous outer bound (wider, but sound, since naïve interval evaluation over-approximates a monotone function) and for p-box inputs would use the Fréchet combination for the unconditioned remainder. We emphasise that this hybrid is proposed, not implemented or evaluated here; it is future work, and the method reported in this paper does not degrade exactness anywhere.

Imprecise propagation. Interval results are exact at machine precision; the speed comparison of Section 5.3.1 is scoped to one-shot queries. Distributional (p-box) bounds are sound but their tightness is structure-dependent




Figure 9: Reliability map of the dense decentralised configuration: each facility coloured by the lower bound of its propagated reachability reliability (left; darker = less certainly reachable) and by the width of its reliability band (right; darker = more uncertainty). Hubs (diamonds) are assumed always available. Uncertainty concentrates at island and far-northern facilities; the worst-served facility, Islay Hospital, is bounded at $[0.56, 0.72]$.

— tight for high-reliability or weakly reconvergent systems, conservative for strongly reconvergent systems with broadly uncertain components — and their cost grows quadratically with the discretisation level, which confines full distributional propagation to modest networks at present; for larger networks interval propagation combined with sampling is the practical recourse. The tighter positive-dependence conditioning operator is validated but not proven sound; the proven Fréchet operator is used for every guaranteed claim, and a proof (or refutation) for the tighter operator is an open mathematical question we have verified is not settled in the existing dependence-bound literature.

Structural scope. The method addresses directed acyclic dependency structures. The applied case study’s networks are directed by the hub-and-spoke supply logic of their source design (Section 5.4.1), which is the appropriate

reading of that system; general cyclic or bidirectional dependency networks — feedback loops, return flows — are outside the present scope, and their principled treatment (for example by unrolling or by strongly-connected-component condensation) is future work.

6. Conclusions

An exact reliability analysis method for directed acyclic process networks with reconvergent dependency structures, the Information Propagation Algorithm (IPA), has been presented. The method computes source-to-node reliability by propagating reachability probabilities in topological order, resolving dependent incoming signals by conditioning on minimal separating sets of shared ancestors, and replacing resolved fork-join substructures with statistically equivalent supernodes. The revision establishes the method’s formal position: it is a specialisation of cutset conditioning to source-to-node reachability, its per-instance cost is exactly $\sum_d 2^{|C_d|}$ over the diamonds resolved — computable before enumeration — and an independent-substructure factorisation reduces fan-in conditioning from exponential to linear where independence permits.

Exactness was verified against an independent exact method, a reduced-ordered binary decision diagram with dynamic variable reordering, across 129 random and mutated networks, six topological families, and real infrastructure networks, with worst observed disagreement at floating-point precision. The same comparison shows that the method’s cost is governed by the same structural width parameter as the diagram’s size, with neither method dominating: for exact point-valued reliability the method is competitive with, not superior to, a well-ordered decision diagram. Its distinguishing capability is the native propagation of imprecise component reliabilities: interval-valued inputs yield the exact reliability range at machine precision — and, for one-shot queries, faster than the decision-diagram route to the same interval on every family tested — while probability-box inputs yield guaranteed distributional bounds via a conditioning operator grounded in the Fréchet–Hoeffding inequalities, enabling certified bounds on requirement-violation probabilities from a single propagation, a result no point-valued exact method produces analytically.

An applied case study of a medical drone logistics network for Scotland, with every input traced to the source design study and one clearly flagged extension, demonstrated the engineering use of the method: propagated re-

liability bounds identify the facilities whose delivery reachability is most uncertain under honest input uncertainty, and a sweep of a real redundancy design parameter located the practical boundary of exact analysis on the network — within which the method completed full exact propagation in under 25 seconds at a redundancy level where the decision-diagram alternative did not complete within a practical budget. Exact reliability analysis of large acyclic networks is feasible whenever dependency-inducing substructures remain sufficiently localised, and the boundary of feasibility can be measured and expressed in design terms rather than assumed.

Future work includes a soundness proof for the tighter positive-dependence conditioning operator, the depth-limited hybrid scheme for networks beyond the exact range, extension to multi-state components and explicitly modelled common-cause failures, and principled treatment of cyclic dependency structures.

Acknowledgements

This work was supported by the Nuclear Decommissioning Authority (NDA) in conjunction with the National Nuclear Laboratory (NNL) through the NDA Bursary Scheme (project number NNL/UA/118).

Data and Software Availability

The implementation of the Information Propagation Algorithm (IPA), together with the benchmark datasets and scripts required to reproduce the case studies, is available at:

<https://github.com/Temi-Tory/InformationPropagationAnalysis.jl>.

References

- Au, S.-K., Beck, J.L., 2001. Estimation of small failure probabilities in high dimensions by subset simulation. *Probabilistic Engineering Mechanics* 16 (4), 263–277.
- Billinton, R., Allan, R.N., 1992. Network modelling and evaluation of simple systems. In: *Reliability Evaluation of Engineering Systems: Concepts and Techniques*. Springer US.
- Brace, K.S., Rudell, R.L., Bryant, R.E., 1990. Efficient implementation of a BDD package. In: *Proceedings of the 27th ACM/IEEE Design Automation Conference*, 40–45.

- Bryant, R.E., 1986. Graph-based algorithms for Boolean function manipulation. *IEEE Transactions on Computers* C-35 (8), 677–691.
- Bryant, R.E., Meinel, C., 2002. Ordered binary decision diagrams. In: *Logic Synthesis and Verification*. Springer.
- Čepin, M., 2011. Event tree analysis. In: *Assessment of Power System Reliability: Methods and Applications*. Springer London.
- Distefano, S., Puliafito, A., 2007. Dynamic reliability block diagrams vs dynamic fault trees. In: *2007 Proceedings — Annual Reliability and Maintainability Symposium*, 71–76.
- Ebeling, C.E., 2019. *An Introduction to Reliability and Maintainability Engineering*, third ed. Waveland Press.
- Ferson, S., Kreinovich, V., Ginzburg, L., Myers, D.S., Sentz, K., 2003. *Constructing Probability Boxes and Dempster–Shafer Structures*. Sandia National Laboratories, SAND2002-4015.
- Fishman, G.S., 1986. A comparison of four Monte Carlo methods for estimating the probability of s–t connectedness. *IEEE Transactions on Reliability* 35 (2), 145–155.
- George-Williams, H., Santhosh, T.V., Patelli, E., 2021. *Simulation Methods for the Analysis of Complex Systems*. Springer.
- Hui, K.-P., 2007. Monte Carlo network reliability ranking estimation. *IEEE Transactions on Reliability* 56 (1), 50–57.
- Jones, D., Filippi, G., Basu, S., Parsonage, B., Patelli, E., Maddock, C., Vasile, M., Fossati, M., 2025. Conceptual design of a medical drone logistics network for Scotland. *International Journal of Logistics Research and Applications* (submitted). [*Author note: update publication status at proof stage.*]
- Kim, Y., Kang, W.-H., 2013. Network reliability analysis of complex systems using a non-simulation-based method. *Reliability Engineering & System Safety* 110, 80–88.
- Kumar, T.B., Sekhar, O.C., Ramamoorthy, M., 2018. Composite power system reliability evaluation using modified minimal cut set approach. *Alexandria Engineering Journal* 57 (4), 2521–2528.
- Lauritzen, S.L., Spiegelhalter, D.J., 1988. Local computations with probabilities on graphical structures and their application to expert systems. *Journal of the Royal Statistical Society, Series B* 50 (2), 157–224.
- Li, H., Zhao, Q., 2005. A cut/tie set method for reliability evaluation of control systems. In: *Proceedings of the 2005 American Control Conference*, 1048–1053.
- Liu, W., Li, J., 2012. An improved cut-based recursive decomposition algorithm for reliability analysis of networks. *Earthquake Engineering and Engineering Vibration* 11 (1), 1–10.

- Mosleh, A., 1993. Estimation of parameters of common cause failure models. In: *Reliability and Safety Assessment of Dynamic Process Systems*. Springer.
- Murphy, K., Weiss, Y., Jordan, M., 2013. Loopy belief propagation for approximate inference: an empirical study. In: *Proceedings of the 15th Conference on Uncertainty in Artificial Intelligence*.
- Pearl, J., 1988. *Probabilistic Reasoning in Intelligent Systems: Networks of Plausible Inference*. Morgan Kaufmann.
- Rausand, M., 2014. *Reliability quantification*. In: *Reliability of Safety-Critical Systems*. John Wiley & Sons.
- Rauzy, A., 1993. New algorithms for fault trees analysis. *Reliability Engineering & System Safety* 40 (3), 203–211.
- Sherwin, D.J., Bossche, A., 1993. Fault trees, event trees and success trees. In: *The Reliability, Availability and Productiveness of Systems*. Springer Netherlands.
- Tong, Y., Tien, I., 2019. Probability propagation method for reliability assessment of acyclic directed networks. *ASCE-ASME Journal of Risk and Uncertainty in Engineering Systems, Part A: Civil Engineering* 5 (3), 04019011.
- Williamson, R.C., Downs, T., 1990. Probabilistic arithmetic I: numerical methods for calculating convolutions and dependency bounds. *International Journal of Approximate Reasoning* 4 (2), 89–158.
- Xing, L., Amari, S.V., 2008. Fault tree analysis. In: *Handbook of Performability Engineering*. Springer London.
- Yedidia, J.S., Freeman, W.T., Weiss, Y., 2003. Understanding belief propagation and its generalizations. In: *Exploring Artificial Intelligence in the New Millennium*. Morgan Kaufmann.